

## Pipeline Hazards

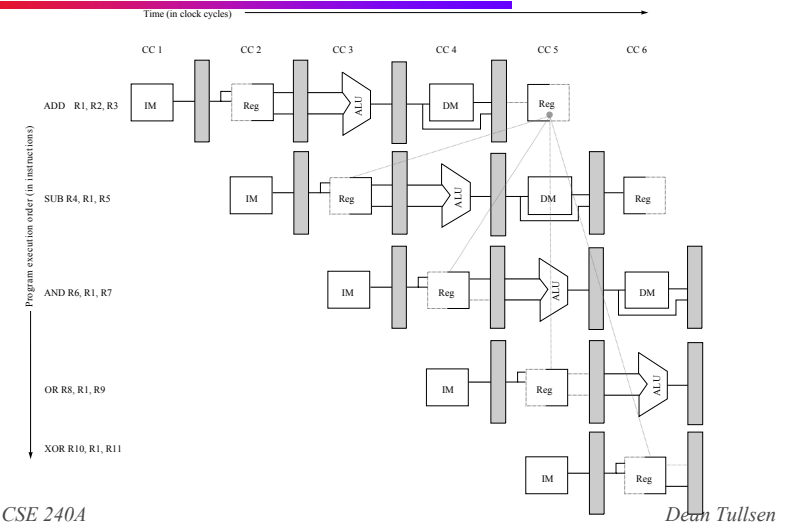
or

*Danger! Danger! Danger!*

CSE 240A

Dean Tullsen

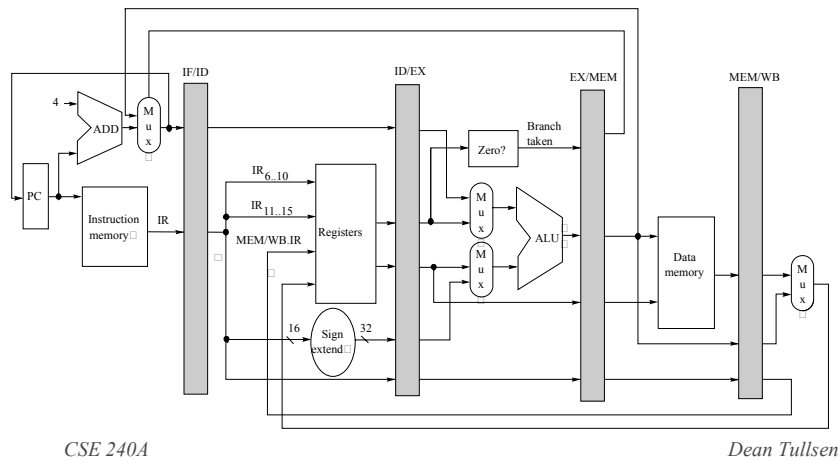
## Data Hazards



CSE 240A

## Data Hazard

lw R8, 10000(R3)    add R6, R2, R1    addi R3, R1, #35



CSE 240A

## Data Dependence

- Data hazards are caused by data dependences
- Data dependences, and thus data hazards, come in 3 flavors (not all of which apply to this pipeline).
  - RAW (read-after-write)
  - WAW (write-after-write)
  - WAR (write-after-read)

CSE 240A


Dean Tullsen

## RAW Hazard

- later instruction tries to read an operand before earlier instruction writes it

- The dependence
  - add **R1**, R2, R3
  - sub R5, **R1**, R4

- The hazard



- RAW hazard is extremely common

CSE 240A


Dean Tullsen

## WAW Hazard

- later instruction tries to write an operand before earlier instruction writes it

- The dependence
  - add **R1**, R2, R3
  - sub **R1**, R2, R4

- The hazard



- WAW hazard possible in a reasonable pipeline, but not in the very simple pipeline we're assuming.

CSE 240A

Dean Tullsen

## WAR Hazard

- later instruction tries to write an operand before earlier instruction reads it

- The dependence
  - add R1, **R2**, R3
  - sub **R2**, R5, R4

- The hazard?

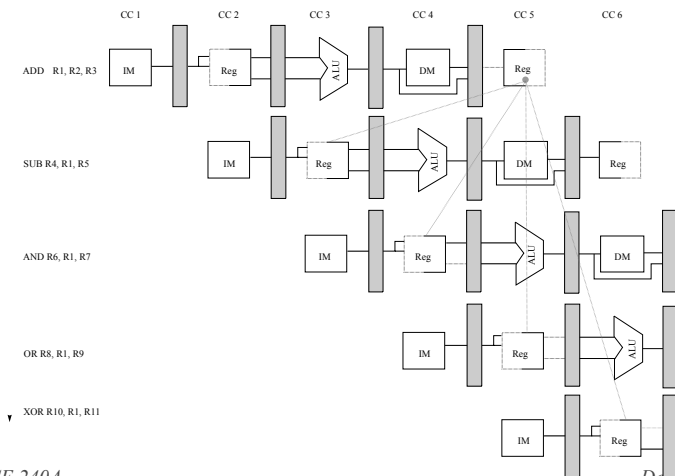


- WAR hazard is uncommon/impossible in a reasonable (in-order) pipeline

CSE 240A

Dean Tullsen

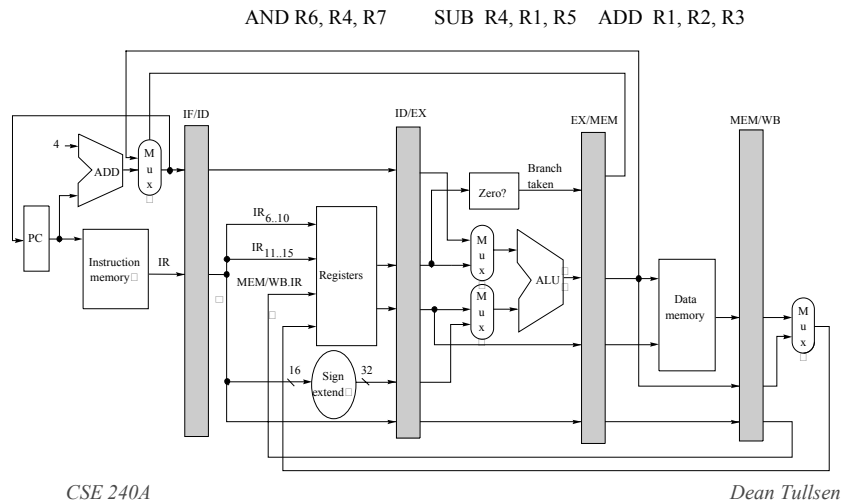
## Dealing with Data Hazards through Forwarding



CSE 240A

Dean Tullsen

## Dealing with Data Hazards through Forwarding



## Forwarding Options

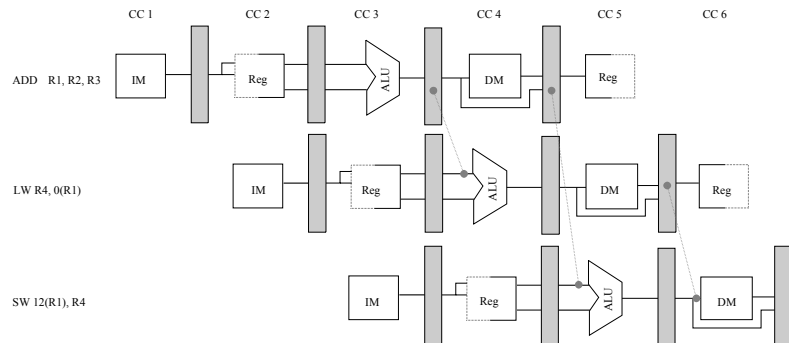
- ADD -> ADD
- ADD -> LW
- ADD -> SW (2 operands)
- LW -> ADD
- LW -> LW
- LW -> SW (2 operands)

(I'm letting ADD stand in for all ALU operations)

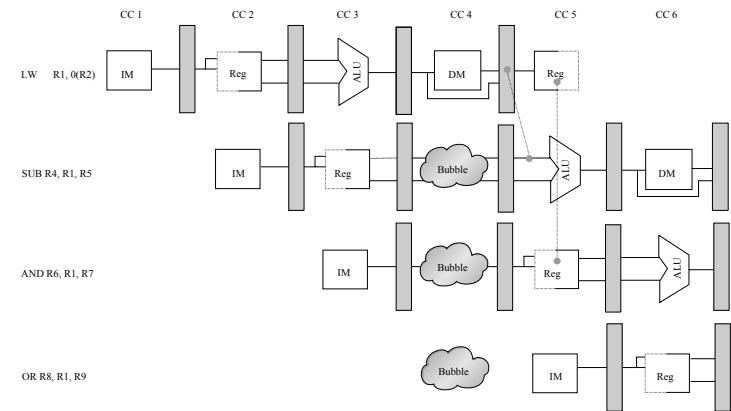
CSE 240A

Dean Tullsen

## More Forwarding



## Forwarding and Stalling



## Example

```
ADD R1, R2, R3
SW R1, 1000(R2)
LW R7, 2000(R2)
ADD R5, R7, R1
LW R8, 2004(R2)
SW R7, 2008(R8)
ADD R8, R8, R2
LW R9, 1012(R8)
SW R9, 1016(R8)
```

CSE 240A

Dean Tullsen

## Avoiding Pipeline Stalls

```
lw R1, 1000(R2)
lw R3, 2000(R2)
add R4, R1, R3
lw R1, 3000(R2)
add R6, R4, R1
sw R6, 1000(R2)
```

- this is a compiler technique called *instruction scheduling*.

CSE 240A

Dean Tullsen


## How big a problem are these pipeline stalls?

- 13% of the loads in FP programs
- 25% of the loads in integer programs

CSE 240A

Dean Tullsen

## Detecting ALU Input Hazards



CSE 240A

Dean Tullsen

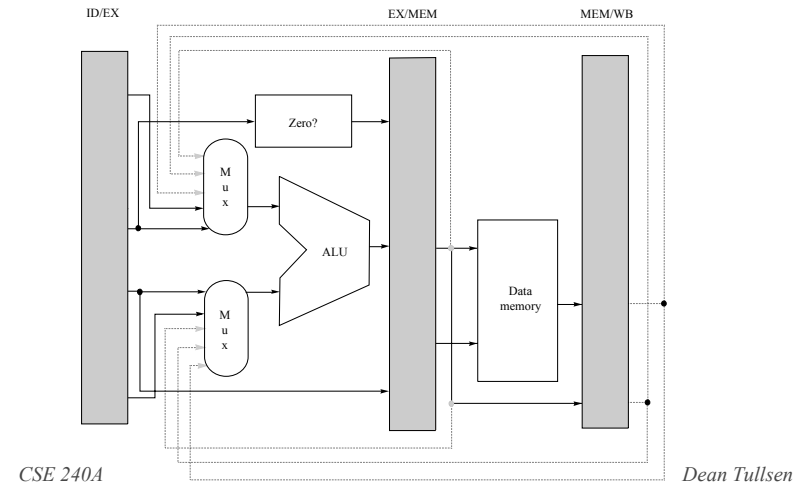
## Inserting Bubbles

- Set all control values in the EX/MEM register to safe values (equivalent to a nop)
- Keep same values in the ID/EX register and IF/ID register
- Keep PC from incrementing

CSE 240A

Dean Tullsen

## Adding Datapaths




CSE 240A

Dean Tullsen

## Control Hazards


- Instructions are not only dependent on instructions that produce their operands, but also on all previous control flow (branch, jump) instructions that lead to that instruction.



CSE 240A

Dean Tullsen

## Branch Hazards



CSE 240A

Dean Tullsen


## Branch Stall Impact

- If CPI = 1, 30% branch, Stall 3 cycles => new CPI = 1.9!
- Two part solution:
  - Determine branch taken or not sooner, AND
  - Compute taken branch address earlier
- (limited MIPS) branch tests if register = 0 or  $\neq 0$
- MIPS Solution:
  - Move Zero test to ID/RF stage
  - Adder to calculate new PC in ID/RF stage
  - 1 clock cycle penalty for branch versus 3

CSE 240A

Dean Tullsen


## New Datapath



CSE 240A

Dean Tullsen

## Branch Hazards



CSE 240A

Dean Tullsen

## What We Know About Branches

- more conditional branches than unconditional
- more forward than backward
- 67% of branches taken
- backward branches taken 80%

CSE 240A

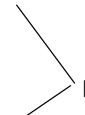
Dean Tullsen

## Four Branch Hazard Alternatives

- #1: Stall until branch direction is clear
- #2: Predict Branch Not Taken
  - Execute successor instructions in sequence
  - “Squash” instructions in pipeline if branch actually taken
  - Advantage of late pipeline state update
  - 33% MIPS branches not taken on average
  - PC+4 already calculated, so use it to get next instruction
- #3: Predict Branch Taken
  - 67% MIPS branches taken on average
  - *But haven't calculated branch target address in this MIPS architecture*
    - MIPS still incurs 1 cycle branch penalty
    - Other machines: branch target known before outcome

## Four Branch Hazard Alternatives

- #4: Delayed Branch
    - Define branch to take place *AFTER* a following instruction
- branch instruction  
sequential successor<sub>1</sub>  
sequential successor<sub>2</sub>  
.....  
sequential successor<sub>n</sub>  
branch target if taken



**Branch delay of length  $n$**
- 1 slot delay allows proper decision and branch target address in 5 stage pipeline
  - MIPS uses this

## Delayed Branch

- Where to get instructions to fill branch delay slot?
  - Before branch instruction
  - From the target address: only valuable when branch taken
  - From fall through: only valuable when branch not taken
  - Cancelling branches allow more slots to be filled
- Compiler effectiveness for single branch delay slot:
  - Fills about 60% of branch delay slots
  - About 80% of instructions executed in branch delay slots useful in computation
  - About 50% (60% x 80%) of slots usefully filled

## Key Points

- Hard to keep the pipeline completely full
- Data Hazards require dependent instructions to wait for the producer instruction
  - Most of the problem handled with forwarding (bypassing)
  - Sometimes stall still required (especially in modern processors)
- Control hazards require control-dependent (post-branch) instructions to wait for the branch to be resolved